

ThetaTrap

A bounded MCP-native AI options agent

Qwen investigates earnings risk and calls Alpaca's official MCP tools. Deterministic policy keeps every paper order inside fixed limits.

PAPER TRADING / BASIC INDICATIVE DATA




The thesis was simple. Safe autonomy was not.

**A YouTube lesson made me ask:
can an agent collect elevated
earnings premium without betting
on direction?**

The hard part was not drawing an iron condor. It was proving the event, liquidity, order state, and exit were safe after timeouts and restarts.

01	Verify	Use first-party earnings timing and require listed expirations.
02	Structure	Sell near-term premium with farther-out wings that cap loss.
03	Decide	Let Qwen inspect evidence and approve or veto the candidate.
04	Exit	Reconcile through Alpaca MCP and target a next-morning close.

The LLM makes the judgment; policy signs the order.



Hard boundary: Qwen cannot change the symbol, strikes, quantity, loss cap, account, or authorization date.

Autonomy stops where deterministic risk begins.

Every entry must pass account identity, market schedule, event, data quality, liquidity, pricing, one-shot authorization and broker-flat gates.

\$500

Maximum loss

Also capped at 0.5% of account equity.

1

Contract

One position at a time; no overlapping exposure.

1

Submission / date

A broker submission consumes that day's authorization.

The release shows evidence, including why it refused risk.

The submission reports broker-reconciled paper results only. It does not turn replay output, simulated fills, or a no-trade decision into a profitability claim.

181

Automated tests

The integrated release covers public mode and read-only smoke.

7

Eligible events

Four Sep 1 and three Sep 2 after-close candidates.

2

Explicit exclusions

Ambiguous timing and unavailable weekly expirations fail closed.

Final equity and P&L are read from Alpaca and displayed on the live dashboard after reconciliation.

Next: better evidence, with the same hard boundaries.



ThetaTrap's goal is not unconstrained intelligence. It is accountable autonomy.